

# Examensarbete

## “Exploring Lightweight Kubernetes Clusters: Installation, Trends, and Implications”

<b>Namn</b>	Josip Jovanovic, Javad Ein Moghassemi
<b>Utbildning</b>	Devops Developer
<b>Uppdragsgivare</b>	EC Utbildning, Fabled AB
<b>Handledare/ examinator</b>	Shahla Yagoutkar, Viktor Svensson
<b>Datum</b>	16 May 2023

## **Table of Content:**

- 1. Abstract
- 2 Introduction
  - 2.1 Background
  - 2.2 Purpose
  - 2.3 Problem Formulation
  - 2.4 Limitations
- 3. Method
- 4. Theory
- 5. Results
  - 5.1 Interview
  - 5.2. Installation And Tutorial
- 6. Discussion
- 7. Conclusion
- 8. Attachments
- 9. Reference List

## **1. Abstract:**

This study investigates the setup and maintenance of lightweight Kubernetes (K8s) clusters that have monitoring and CI/CD pipelines capabilities. Alternatively, the study is a limited survey of the K8s ecosystem, focusing on programs such as K3D (a lightweight K8s distributor), Grafana and Prometheus (for monitoring features), and Tekton (for CI/CD pipelines that are native to K8s). A secondary goal of this study is to facilitate new trainees with a tutorial that includes the installation process of this K8s bundle, and provide a brief description about each of its components. This report is written by two DevOps students at EC Utbildning who have combined their education with their LIA experience in order to equip individuals with the necessary knowledge and skills to leverage K8s effectively.

## 2. Introduction

### 2.1 Background:

Kubernetes (K8s), is the first graduate project (made available as a fully open source software) of the Cloud Native Computing Foundation (CNCF), and it has been majorly adopted to help manage swarms of containers within the IT industry since its initial release in September 2014 (Kubernetes, 2023). In its *State of Enterprise Open Source 2022* report, Redhat details that 70 percent of the 1,296 companies participating in its survey use K8s (Enterprise, 2023). Another report published the same year by Datadog titled *9 Insights on Real-World Container Usage*, describes that “nearly half of all organizations using containers run K8s to deploy and manage some of those containers” (CNCF, 2022). It must also be noted that the widespread and robust adaptation of K8s is partly due to the variety of K8s flavors that are offered by different distributors.

Some distributors offer basic K8s clusters (what is better known as “vanilla” K8s), whereas cloud computing platforms, such as Azure and AWS, offer clusters that are managed internally. There are also distributors that offer lightweight clusters with less dependencies. According to CNCF’s annual survey, in 2021, 30 percent of the participants used *Minikube* and 22 percent used *K3s*, which are both lightweight K8s distributors (CNCF, 2021).

To underline the significance of lightweight clusters, it is imperative to consider various scenarios where they offer a superior solution for computing needs. These scenarios include situations where access to abundant hardware power is limited, cost optimization is a priority, greater control over troubleshooting the clusters is desired, specific data storage locations are required, or the flexibility to relocate data to different physical locations is essential. In such cases, relying solely on cloud-based services may not be the optimal choice. Alternatively, the utilization of customized lightweight clusters, equipped with a curated selection of open-source software, proves to be highly advantageous compared to alternative options.

Moreover, according to CNCF’s surveys, training new K8s users is a major obstacle within the industry. Comprehensive tutorials play a pivotal role in enhancing the learning process. While numerous online tutorials are available that primarily focus on individual programs within the Kubernetes ecosystem, the availability of tutorials that investigate a collection of programs is relatively limited. This lack of such comprehensive resources has served as inspiration for producing this report.

## **2.2 Purpose**

The objective of this study is to provide a comprehensive guide on establishing a lightweight and high-performing Kubernetes cluster with a complete feature set for CI/CD pipelines and observability. The installation process will also be documented to create a tutorial that can be used for training purposes. In addition to the cluster, other major phases of this report include deploying an application with a database to the cluster, and using K8s' exclusive package manager known as Helm, for the deployments.

## **2.3 Problem Formulation**

This study investigates a collection of open-source programs that can be integrated into K8s clusters. It provides comprehensive information about the installation procedures, along with a brief description of each component. Furthermore, the study addresses prevalent challenges that users may encounter when working with these technologies.

The following topics will be investigated in this research:

- 1) Provide a concise overview of a few essential features of K8s components and create a tutorial based on it.
- 2) Determine which tools and software within the Kubernetes ecosystem can facilitate the deployment of clusters with CI/CD pipelines and observability features, particularly in scenarios where minimizing hardware usage is a priority.
- 3) Create a set of instructions for installing and troubleshooting the selected K8s bundle.

## **2.4 Limitations**

This study acknowledges certain limitations concerning the depth of coverage regarding K8s' architecture. Due to the extensive number of components involved in K8s, it is deemed satisfactory to provide a fundamental understanding at a basic level, rather than delving into intricate details. The primary objective is to demonstrate the capabilities of the examined software to new students, offering a beginner-level perspective.

Additionally, it should be noted that the installation of prerequisites for the project will not be addressed within the scope of this study. While the installation of prerequisites is essential for the successful implementation of the project, it is beyond the immediate focus of this research.

### **3. Method**

The primary objective of this study is to examine a specific K8s bundle that demonstrates minimal hardware resource consumption. The outcome of this investigation will result in the development of a replicable model that is accessible for first-time users of Kubernetes. In order to ensure clarity and enhance readability, this research relies extensively on referencing articles, surveys, and official documentation provided by vendors. Additionally, insights derived from an interview conducted with the authors' LIA supervisor, Robin Morero, have been incorporated into the text.

There are other tools and softwares that need to be pre installed when following this tutorial. Notably, the prerequisites for the project include WSL 2 (Windows Subsystem for Linux 2), Docker Desktop, and access to a Linux CLI. The programs in the K8s bundle will be installed in this order: K3D, Grafana, Tekton, Helm, and MongoDB

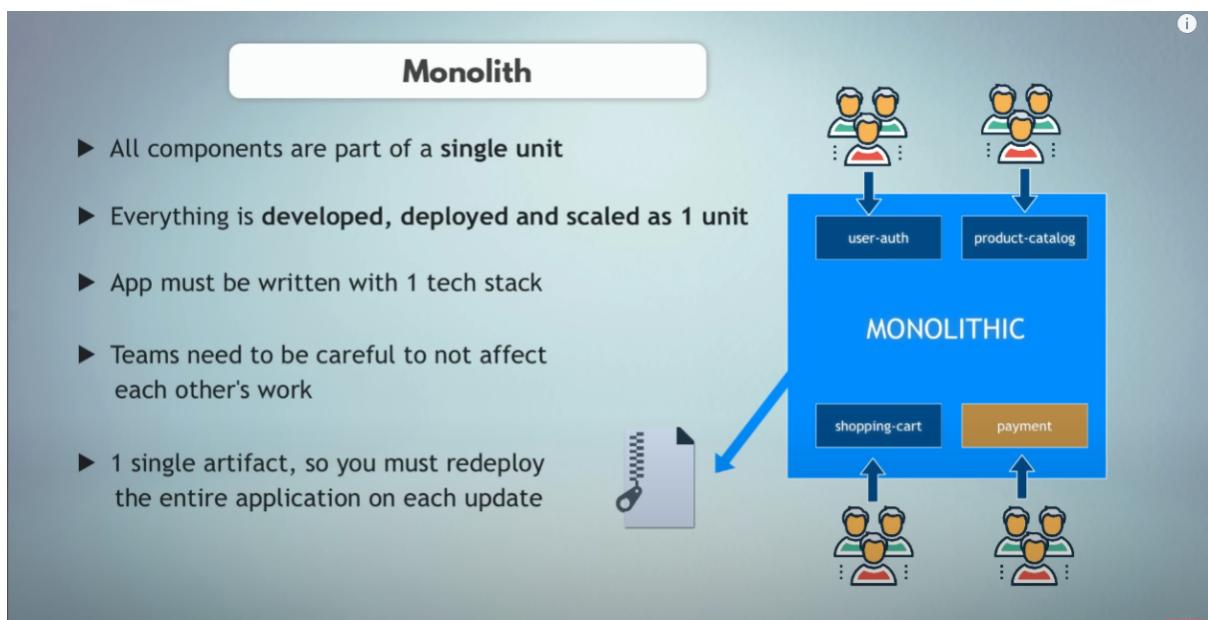
## **4. Theory**

The softwares that are utilized in this study will be explained in the order provided below:

- 1) Monolithic Architecture VS. Microservices (a brief history)
- 2) K3D
- 3) Monitoring
- 4) CI/CD
- 5) Helm
- 6) MongoDB

## Monolith Architecture:

The main characteristic of a Monolithic architecture is that all the different components of a software artifact are considered as a single unit or “stack”. By building software products with a single code-base, and limiting the deployment configurations to a single file, this type of architecture allows for easier development, debugging, and deployment. However, this approach also has some shortcomings. For example, if developers need to make small updates or changes to an artifact, they ought to interact with the entire source code of that artifact, which in turn makes small updates more time consuming and more costly. Furthermore, monolithic software is harder to scale.

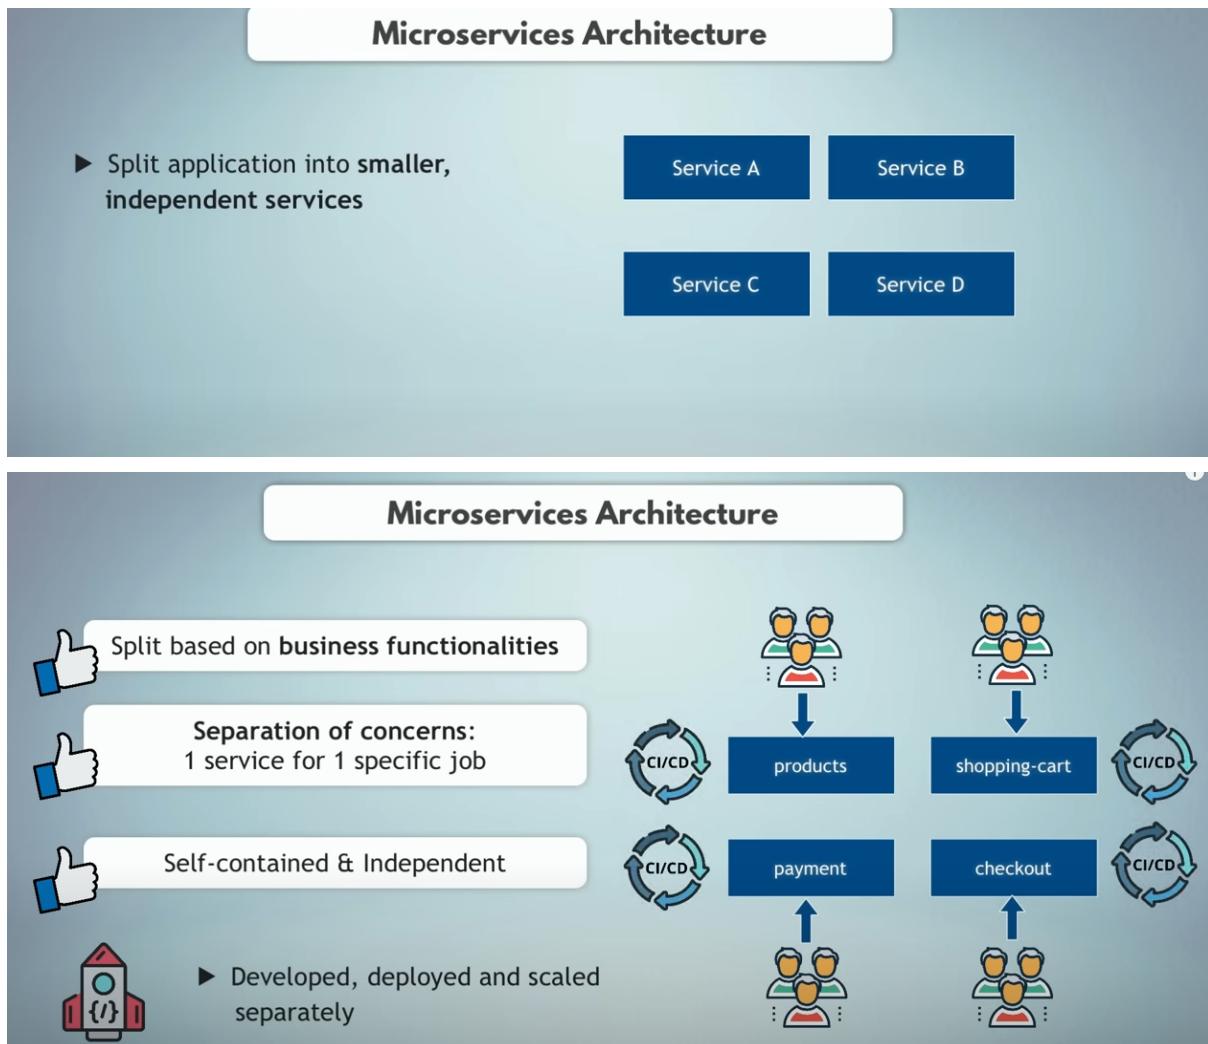


(TechWorld with Nana. 2022)

## Microservices Architecture:

The Microservice architecture divides the different components of an artifact into smaller, independent services, giving rise to micro applications. Since microservices are self-contained and independent, it means that they can be built and deployed separately .

"Monoliths are not always negative. For a simple prototype, or a single team, it may be more effective (Robin Morero, 2023)"



(TechWorld with Nana. 2022)

## **K3D:**

“K3D is a lightweight wrapper to run K3S (Rancher Lab’s minimal Kubernetes distribution) in docker”(K3D, 2022).

K3S is a highly available, certified Kubernetes distribution designed for production workloads in unattended, resource-constrained, remote locations or inside IoT appliances. We will be using K3D for our cluster, for its multi-node capabilities.

## Monitoring (Prometheus and Grafana):

### Prometheus

“Prometheus is an open-source systems monitoring and alerting toolkit originally built at SoundCloud. Since its inception in 2012, many companies and organizations have adopted Prometheus, and the project has a very active developer and user community. It is now a standalone open source project and maintained independently of any company. Prometheus collects and stores its metrics as time series data, i.e. metrics information is stored with the timestamp at which it was recorded, alongside optional key-value pairs called labels”(Prometheus, 2023).

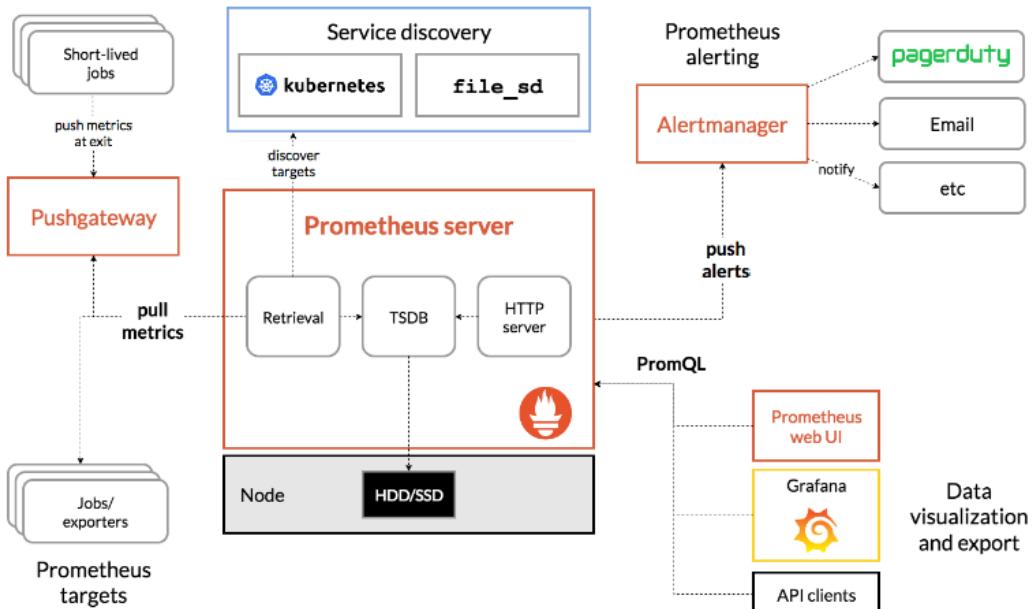
Prometheus uses a language called PromQL for its queries.

This is very useful for our clusters well-being, this way we can monitor our pods and services.

We will install it to our cluster using helm charts.

### Architecture

This diagram illustrates the architecture of Prometheus and some of its ecosystem components:



## Grafana

“Grafana is an [open source](#) interactive data-visualization platform, developed by [Grafana Labs](#), which allows users to see their data via charts and graphs that are unified into one dashboard (or multiple dashboards!) for easier interpretation and understanding” (Grafana 2022).

Grafana goes hand in hand with Prometheus. Grafana can be connected to Prometheus by adding Prometheus as a data source in Grafana. This allows Grafana to retrieve metrics data from Prometheus and display it in dashboards. We will also install grafana through a helm chart.



## Tekton:

“Tekton is a cloud-native solution for building CI/CD pipelines. It consists of Tekton Pipelines, which provides the building blocks, and of supporting components, such as Tekton CLI and Tekton Catalog, that make Tekton a complete ecosystem” (Tekton, 2018).

We decided that we are trying out Tekton to see if this matches what Robin says.

In our practical part we will create a simple pipeline using tekton that gets the information from a github repo and deploys the pipe.

We will also install Tekton Dashboard for its GUI capabilities.

We will install it to our cluster using helm charts.

```
+ gcloud auth activate-service-account `--key-file=/secret/release.json`
+ Added service account credentials for: [release-right-memo@tekton-releases.iam.gserviceaccount.com]
+ gcloud auth configure-docker
Adding credentials for all GCR repositories.
WARNING: A long list of credential helpers may cause delays running 'docker build'. We recommend passing the registry name to:
After update, the following will be written to your Docker config file
located at [/tekton/home/.docker/config.json]:
credHelpers: [
    "gcr.io": "gcloud",
    "us.gcr.io": "gcloud",
    "europe.gcr.io": "gcloud",
    "asia.gcr.io": "gcloud",
    "staging-e85.gcr.io": "gcloud",
    "marketplace.gcr.io": "gcloud"
]

Do you want to continue (Y/n)?
Docker configuration file updated.
+ date
+ export 'SOURCE_DATE_EPOCH=1611112427'
+ cd /workspace/go/src/github.com/tektoncd/dashboard
+ sed -i s/devel/v20210128-00551e3b76/g /workspace/go/src/github.com/tektoncd/dashboard/base/300-deployment.yaml
+ sed -i s/devel/v20210128-00551e3b76/g /workspace/go/src/github.com/tektoncd/dashboard/base/300-service.yaml
+ which ko
+ /usr/local/bin/ko
+ ko version
2021/01/20 03:13:47 NOTICE!
-----
Please install ko from https://github.com/google/ko.
For more information see:
https://github.com/google/ko/issues/258
...
0/0
+ customize version
{Version:unknown GitCommit:$Format:$H$ BuildDate:1970-01-01T00:00:00Z GoOs:linux GoArch:amd64}
```

## **Helm:**

Helm is a package manager for K8s and it can be used to “define, install, and upgrade K8s applications” (Helm). This package manager reduces redundancy through its “Helm Charts” feature. This feature allows to bundle all the yaml files that are required by K8s for deployment (Secret, Configmap, Deployment.yaml, etc.) into a single chart. By utilizing “Helm Charts”, developers can avoid applying yaml files one by one.

**What is Helm?**

**Helm Charts**

- Bundle of YAML Files
- Create your own Helm Charts with Helm
- Push them to Helm Repository
- Download and use existing ones

## **MongoDB**

MongoDB is a popular open-source document-oriented NoSQL database. It uses a flexible and dynamic schema which allows for easy and efficient data modeling, storage, and retrieval. Rather than using tables and rows as in traditional relational databases, MongoDB stores data as JSON-like documents with dynamic schemas, allowing for more flexibility and scalability.

MongoDB is also known for its high performance, scalability, and availability. It can handle large amounts of data, making it well-suited for applications that require high levels of scalability and performance. Additionally, MongoDB provides built-in replication and sharding, which enables data to be distributed across multiple servers for improved availability and scalability.

Other benefits of using MongoDB include its support for a variety of programming languages, ease of use, and rich querying capabilities. It is also highly flexible and can be used for a variety of use cases, including web and mobile applications, content management systems, and analytics platforms.

## MongoDB Express

MongoDB Express is a web-based administrative interface for MongoDB. It allows users to manage MongoDB databases, collections, and documents using a graphical user interface rather than the command-line interface provided by the MongoDB shell. MongoDB Express is often used as a companion tool to MongoDB, providing a simple and intuitive way for developers and administrators to interact with their MongoDB data. It is an open-source project and can be downloaded and installed for free. We will install both with helm.

The screenshot shows the MongoDB Express interface for the 'test' database. The top navigation bar includes a status icon, the title 'Mongo Express', and a dropdown for the 'Database: test'. Below the header, the main content area is titled 'Viewing Database: test'.

**Collections:** This section lists the collections in the database. Each collection is represented by a row in a table with three columns: 'View' (green button), 'Export' (orange button), and '[JSON]' (orange button). To the right of the table, there is a search bar labeled 'Collection Name' and a blue button labeled '+ Create collection'. Below the table, there are six collection names: 'CollectionB', 'collectionA', 'fs.chunks', 'fs.files', 'somebuck.chunks', and 'somebuck.files'. Each collection name has a red 'Del' button to its right.

**GridFS Buckets:** This section lists the GridFS buckets. It contains two rows, each with a 'View' button (green) and a bucket name ('fs' and 'somebuck').

**Database Stats:** This section provides various statistics for the database. The table has two columns: the metric and its value. The metrics listed are:

Database Stats	
Collections (incl. system.namespaces)	6
Data Size	2.10 MB
Storage Size	36.6 MB
Avg Obj Size #	110 KB
Objects #	19
Extents #	0
Indexes #	13
Index Size	418 KB

## 5. Results

### 5.1 Interview

We interviewed Robin Morero on 10 of May, 2023. He is our LIA supervisor at Fabled and we got an interesting interview from him. Robin was talking about his company, how he sees his company moving forward, and based on what criteria they choose to work with an open source software.

**Javad: How does Fabled choose which software to use from the K8s ecosystem?**

**Robin:**

“It should be easy to work with, easy to maintain and it should work well in the cluster. That's the 3 main components i think! It should also be easy to upgrade and a good community and fairly stable. We should also be able to replace components when we need to. Keeping it lightweight supports all of that in a way that more heavy software does not! ”

**Josip: Why did you decide to use Tekton, Grafana, K3D, Database?**

**Robin:**

**Tekton**

“Fabled is using Tekton for its lightweight, flexible capabilities, complex pipelines and they also wanted something that runs well natively in k8s. And at the moment Tekton was our best option, and runs fairly well today! ”

**Grafana**

“ ElasticSearch would be an alternative but Grafana had the possibility of keeping all of its components open-source and it's so lightweight. That's a big deal for us, i mean ElasticSearch has Java involved, OpenSearch is the same, uses a lot more resources and takes a lot more maintenance. Grafana stack is low maintenance, you don't have to do a lot with it! ”

## K3S

“Instead of going to Vanilla K8s, we wanted K3S because it's low on resources. That's the main rivalry. It's not a big deal really it's only 5 percent, but since we have a lot of small customers we want a minimal possible server to run it at because it's cheaper. It makes a difference! So it's a matter of cost effectiveness.”

## Database

“There are a lot of Postgres everywhere, and we tried a different flavor of postgres, we tried Cockroach, Yugabyte, we tried different operators for postgres for it to run in k8s. So far none of the solutions has been ideal, we like postgres but we don't like how it manages in k8s. Yugabyte had a problem of being very slow with starting up creating tables, that's what bothered us with that one. Cockroach is missing some features that were missing from postgres.”

**Josip:** “ Do you think mongoDB is a good viable choice database?

**Robin:** ”MongoDB is a safe bet I think! It's been used in a lot of places, fairly easy to manage, it's a nice developer experience if you worked with document databases before. It depends on the use case! The reason we are using postgres is that the standard MYSQL has been missing some features we have been using and it has been scaling worse. Postgres has become the standard I think in terms of Relational DB mess.”

**Javad:** “ What do you believe are the future trends for these features? ”

**Robin:**

“ What we are seeing is a couple of trends, the first one is to manage the complexity, which means you set a separate platform team that will manage that. You use vendors to manage the complexity, you must keep adding stuff into k8s cluster to handle it as well, to get line of costs to keep track of what it is costing you, you get carbon footprint, you get security visibility, compliance visibility and so on. The other approach is to minimize complexity, that means you make it as lightweight as you can. You reduce the maintenance work you have to do, and you make it very easy for developers to do whatever they really need to do in the cluster. And our approach has always been to minimize the complexity, it works better, because when adding things to the cluster the complexity grows. So we try to avoid adding it when not necessary. I Mean no one is going to really complain to you about your k8s

cluster costing too much if it's not really costing a lot. So if you keep it lightweight you're not gonna get complaints in the beginning.

**Josip:** “ If you wouldn't choose k8s today in your structure, what would you use today? ”

**Robin:** “ Nomad is still an option for k8s today. Not that common but it's still there. K8s dominates the market today! ”

**Javad:** “ Will Grafana and Tekton be relevant in the near future? ”

**Robin:** “ Grafana is growing in popularity right now and people are adopting it and will be relevant in the near future.

Tekton might be replaced by other options because it can get complex, let's say you want callbacks from git repo to have a webhook, that could get complicated. ”

## 5.2 The Installation Tutorial

### *Installing Tekton*

Do the following steps in this order.

Go to this git repo: <https://github.com/JOJO840/Tekton-pipe-demo>  
then do a: git clone <https://github.com/JOJO840/Tekton-pipe-demo.git>  
or with the ssh method: git clone git@github.com:JOJO840/Tekton-pipe-demo.git  
Run the bash file ./tekton-cmds.sh, this command will install Tekton pipelines and dashboard and also get the updates from apt library.

```
$ tekton-cmds.sh x
examensProject > Tekton-pipe-demo > $ tekton-cmds.sh
1  #!/bin/bash
2
3  # 1. update the apt manager
4  sudo apt-get update
5
6  # 2. Installing tekton pipelines and CRDs
7  kubectl apply -f https://storage.googleapis.com/tekton-releases/pipeline/latest/release.yaml
8
9  # 3. Installing Tekton Dashboard and CRDs
10 kubectl apply -f https://storage.googleapis.com/tekton-releases/dashboard/latest/release.yaml
11
12 |
```

If you run this cmd 'kubectl get ns' you will now see a new namespace where tekton-pipelines lives.

```
josip@G14:~$ kubectl get ns
NAME          STATUS   AGE
kube-system   Active   26d
default       Active   26d
kube-public   Active   26d
kube-node-lease Active   26d
tekton-pipelines Active   114s
tekton-pipelines-resolvers Active   113s
```

To see the every pod in that namespace type in this cmd

'kubectl get all -n tekton-pipelines'

NAME	READY	STATUS	RESTARTS	AGE		
pod/tekton-pipelines-controller-85767d6dc8-p6gmj	1/1	Running	0	45m		
pod/tekton-pipelines-webhook-54cc86bf59-dqsbk	1/1	Running	0	45m		
pod/tekton-dashboard-7487756644-tdwjw	1/1	Running	0	15m		
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)		
service/tekton-pipelines-controller	ClusterIP	10.43.34.196	<none>	9090/TCP,8008/TCP,8080/TCP		
service/tekton-pipelines-webhook	ClusterIP	10.43.178.73	<none>	9090/TCP,8008/TCP,443/TCP,8080/TCP		
service/tekton-dashboard	ClusterIP	10.43.77.139	<none>	9097/TCP		
NAME	READY	UP-TO-DATE	AVAILABLE	AGE		
deployment.apps/tekton-pipelines-controller	1/1	1	1	45m		
deployment.apps/tekton-pipelines-webhook	1/1	1	1	45m		
deployment.apps/tekton-dashboard	1/1	1	1	20m		
NAME	DESIRED	CURRENT	READY	AGE		
replicaset.apps/tekton-pipelines-controller-85767d6dc8	1	1	1	45m		
replicaset.apps/tekton-pipelines-webhook-54cc86bf59	1	1	1	45m		
replicaset.apps/tekton-dashboard-7487756644	1	1	1	15m		
replicaset.apps/tekton-dashboard-6694db567c	0	0	0	20m		
NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/tekton-pipelines-webhook	Deployment/tekton-pipelines-webhook	4%/100%	1	5	1	45m

To connect to the Tekton dashboard svc write in this cmd:

kubectl port-forward -n <namespace> svc/<svc.name> port number

in our case:

kubectl port-forward -n tekton-pipelines svc/tekton-dashboard 9097

```
Josip@G14:~/examensProject/Tekton-pipe-demo$ kubectl port-forward -n tekton-pipelines svc/tekton-dashboard 9097
Forwarding from 127.0.0.1:9097 -> 9097
Forwarding from [::]:9097 -> 9097
Handling connection for 9097
```

type in localhost:9097 in your browser, we use chrome. This will redirect you to the Tekton dashboard.

The screenshot shows the Tekton Dashboard interface. On the left, there's a sidebar with navigation links: 'Tekton resources' (Pipelines, PipelineRuns, Tasks, ClusterTasks, TaskRuns, CustomRuns), 'Import resources', 'About Tekton', and 'Settings'. The main content area has a title 'About Tekton' with a sub-section 'Environment details'. Under 'Dashboard', it shows 'Namespace: tekton-pipelines' and 'Version: v0.35.0'. Under 'Pipelines', it shows 'Namespace: tekton-pipelines' and 'Version: v0.47.0'. To the right of the main content is a large, stylized red illustration of a cat wearing a magnifying glass, looking at a document.

## Tekton Import Resources:

To make a pipeline in Tekton we will use the 2 files in the github repo pipeline.yaml and tasks.yaml. <https://github.com/JOJO840/Tekton-pipe-demo/tree/main>

README.md	Initial commit	last week
pipeline.yaml	Create pipeline.yaml	last week
tasks.yaml	Create tasks.yaml	last week
tekton-cmds.sh	added changes in tekton bash script	12 minutes ago

Steps for pipeline:

1. For Tekton to import and create resources to your cluster it needs permission to do so. We can get it at Tektons site at <https://tekton.dev/docs/dashboard/tutorial/> and use the copy to clipboard icon and paste in the commands in your cli.

## Grant required permissions

The import must be executed using a ServiceAccount with permissions to create the resources being imported.

For this tutorial we will create a ClusterRole granting permission to create a number of Tekton resources, and a RoleBinding configuring this so that the `default` ServiceAccount in the `tekton-dashboard` namespace can create resources in the `default` namespace.

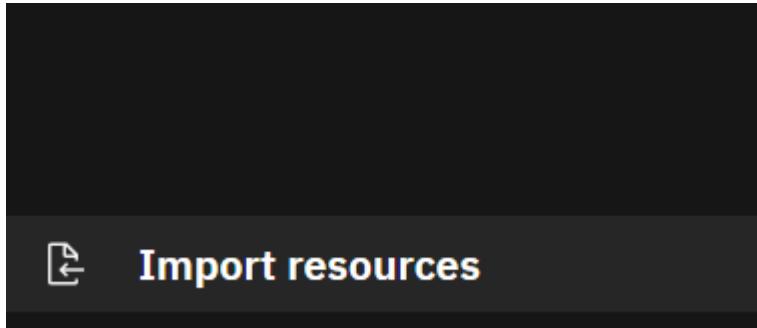
[Copy to clipboard](#)

```
kubectl apply -f - <<EOF
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: tekton-dashboard-tutorial
rules:
  - apiGroups:
    - tekton.dev
    resources:
      - tasks
      - taskruns
      - pipelines
      - pipelineruns
    verbs:
      - get
      - create
EOF
```

```
PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL

● josip@G14:~$ kubectl apply -f - <<EOF
> apiVersion: rbac.authorization.k8s.io/v1
> kind: ClusterRole
> metadata:
>   name: tekton-dashboard-tutorial
>   rules:
>     - apiGroups:
>       - tekton.dev
>         resources:
>           - tasks
>             - taskruns
>             - pipelines
>             - pipelineruns
>           verbs:
>             - get
>             - create
>             - update
>             - patch
>   ---
>   apiVersion: rbac.authorization.k8s.io/v1
>   kind: RoleBinding
>   metadata:
>     name: tekton-dashboard-tutorial
>     namespace: default
>   roleRef:
>     apiGroup: rbac.authorization.k8s.io
>     kind: ClusterRole
>     name: tekton-dashboard-tutorial
>   subjects:
>     - kind: ServiceAccount
>       name: default
>       namespace: tekton-dashboard
> EOF
clusterrole.rbac.authorization.k8s.io/tekton-dashboard-tutorial created
rolebinding.rbac.authorization.k8s.io/tekton-dashboard-tutorial created
○ josip@G14:~$
```

2. Now in the Tekton dashboard Click on Import Resources



3. Put in the git repo and namespace to default like the picture below and click on import.

Tekton Dashboard

**Tekton resources**

- Pipelines
- PipelineRuns
- Tasks
- ClusterTasks
- TaskRuns
- CustomRuns

**Import resources**

Repository URL ⓘ

https://github.com/JOJO840/Tekton-pipe-demo

Repository path (optional) ⓘ

Enter repository path

Revision (optional) ⓘ

Enter revision

Target namespace ⓘ

default

Advanced configuration for the Import PipelineRun

Import

Tekton Dashboard

**Tekton resources**

- Pipelines
- PipelineRuns**
- Tasks
- ClusterTasks
- TaskRuns
- CustomRuns

**import-resources-1683879184205** Last updated 17 seconds ago

Succeeded Tasks Completed: 2 (Failed: 0, Cancelled 0, Skipped: 0)

Step	Status	Duration
fetch-repo	Completed	
clone	Completed	Duration: 1s
import-resources	Completed	

Logs

```
{"level": "info", "ts": 1683879194.2201538, "caller": "git/git.go:176", "msg": "Successful"}, {"level": "info", "ts": 1683879194.2304173, "caller": "git/git.go:215", "msg": "Successful"}
```

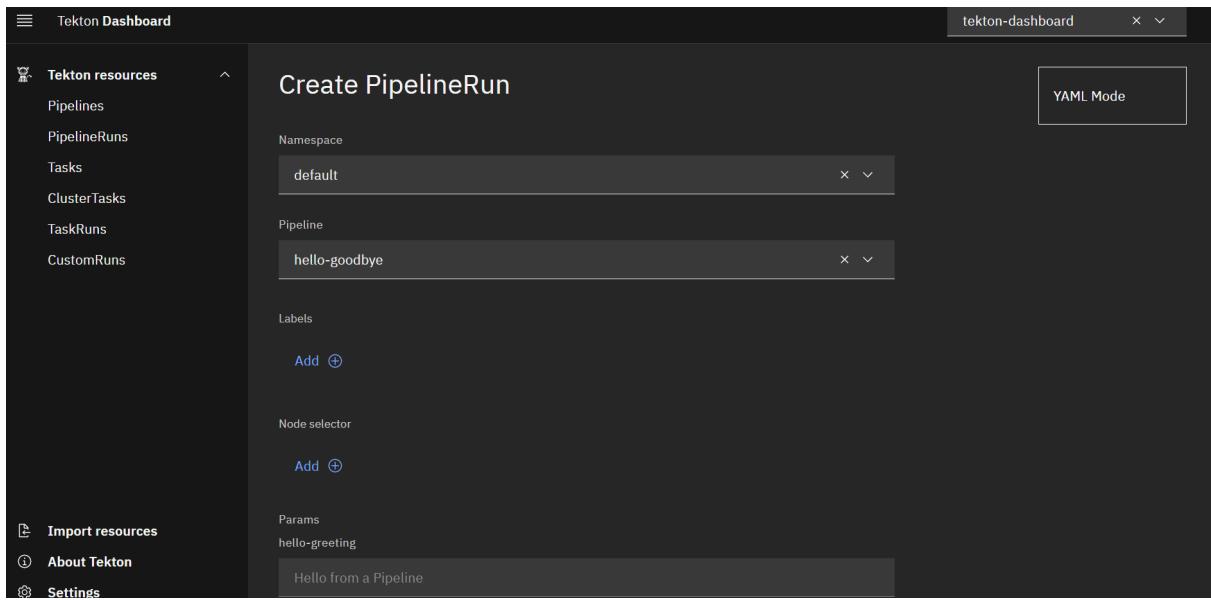
Step completed successfully

you should have two pods up and running in your cluster like this:

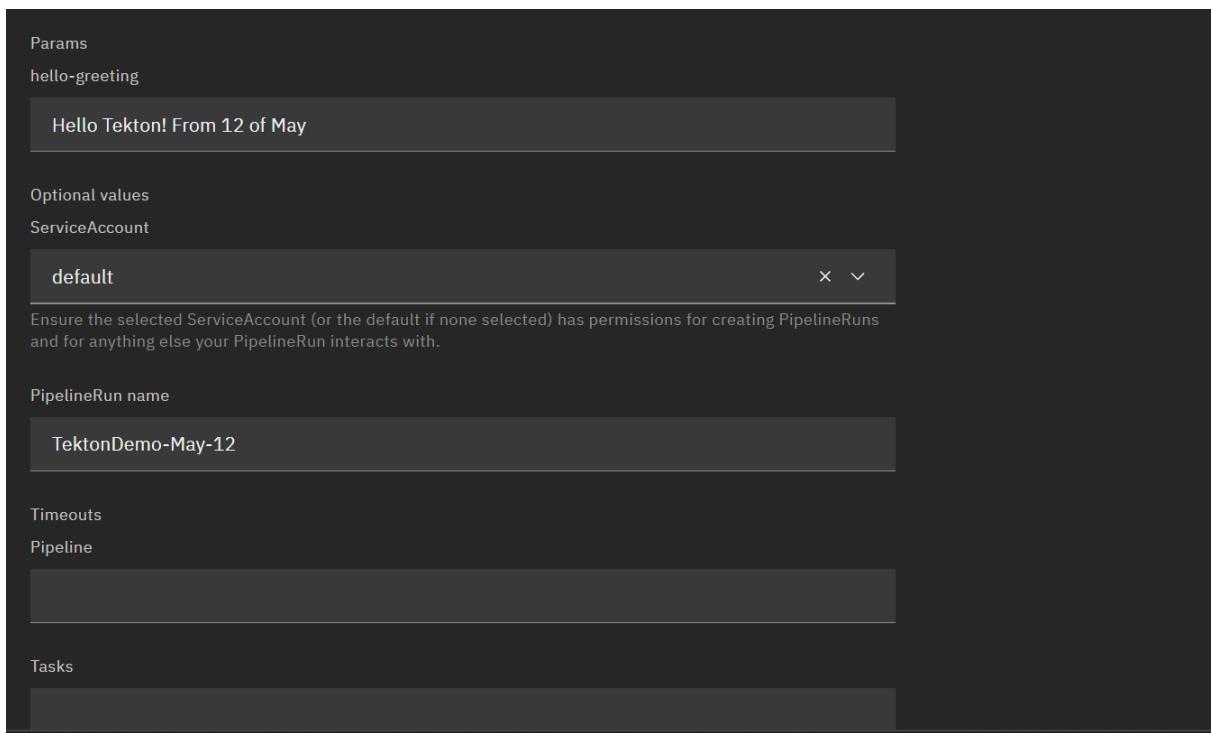
```
● josip@G14:~$ kubectl get pods -n tekton-dashboard
NAME                                         READY   STATUS    RESTARTS   AGE
import-resources-1683879184205-fetch-repo-pod   0/1     Completed   0          107s
import-resources-1683879184205-import-resources-pod   0/1     Completed   0          97s
○ josip@G14:~$
```

## Create a PipelineRun

1. Next let's create a PipelineRun using all of the resources we imported



2. Here you can also choose paragraphs for your pipeline and name it. Then click blue button “Create”



3. Now the pipeline should be green and you will also see some more pods in your cluster.

The screenshot shows the Tekton Dashboard interface. On the left, there's a sidebar with 'Tekton resources' expanded, showing 'Pipelines', 'PipelineRuns', 'Tasks', 'ClusterTasks', 'TaskRuns', and 'CustomRuns'. Under 'PipelineRuns', 'tektondemo-may-12' is selected. The main panel displays the details of this PipelineRun. It shows the status as 'Succeeded' with 'Tasks Completed: 2 (Failed: 0, Cancelled 0, Skipped: 0)'. Below this, three tasks are listed: 'hello' (green checkmark), 'echo' (green checkmark), and 'goodbye' (green checkmark). A 'Parameters' section shows a single entry: 'greeting' with the value 'Hello Tekton! From 12 of May'. The 'Status' and 'Pod' tabs are also visible but contain no data.

4. type in kubectl get pods -A, you will see two new pods in the default namespace now.

```
● josip@G14:~$ kubectl get pods -A
NAMESPACE     NAME                               READY   STATUS    RESTARTS   AGE
kube-system   helm-install-traefik-crd-z6x8s   0/1    Completed  0          27d
kube-system   helm-install-traefik-8dj2r      0/1    Completed  1          27d
tekton-pipelines tekton-dashboard-748775644-tdwjb 1/1    Running   1 (98m ago) 43h
tekton-pipelines tekton-pipelines-controller-85767d6dc8-p6gmj 1/1    Running   1 (98m ago) 43h
kube-system     traefik-66c46d954f-nvlg6      1/1    Running   5 (98m ago) 27d
default        nginx-76d6c9b8c-5bnzj       1/1    Running   5 (98m ago) 27d
kube-system     coredns-597584b69b-qtc4s     1/1    Running   5 (98m ago) 27d
tekton-pipelines-resolvers tekton-pipelines-remote-resolvers-857bc4859-58r2h 1/1    Running   1 (42h ago) 43h
kube-system     metrics-server-5f9f776df5-dxrhr 1/1    Running   6 (98m ago) 27d
kube-system     svc1b-traefik-ddd847c9-tcrmm  2/2    Running   14 (98m ago) 27d
kube-system     svc1b-traefik-ddd847c9-tnbqf  2/2    Running   14 (98m ago) 27d
tekton-pipelines tekton-pipelines-webhook-54cc86bf59-dqsbk 1/1    Running   1 (98m ago) 43h
kube-system     local-path-provisioner-79f67d76f8-dtmkk 1/1    Running   8 (98m ago) 27d
kube-system     svc1b-traefik-ddd847c9-nkkn5  2/2    Running   12 (98m ago) 27d
tekton-dashboard import-resources-1683879184205-fetch-repo-pod 0/1    Completed  0          26m
tekton-dashboard import-resources-1683879184205-import-resources-pod 0/1    Completed  0          26m
default        tektondemo-may-12-hello-pod   0/1    Completed  0          3m51s
default        tektondemo-may-12-goodbye-pod 0/1    Completed  0          3m44s
○ josip@G14:~$
```

## Installing Grafana, Prometheus and Loki

1. Firstly we need to add the repos to Helm for the installation. Run these commands:
2. helm repo add prometheus-community  
<https://prometheus-community.github.io/helm-charts>
3. helm repo add grafana <https://grafana.github.io/helm-charts>
4. helm repo list

```
● josip@G14:~$ helm repo list
NAME      NAMESPACE   REVISION      UPDATED STATUS    CHART          APP VERSION
bitnami   default     1            2023-05-12 12:23:12 1d ago  stable        12.16.0
ingress-nginx  default     1            2023-05-12 12:23:12 1d ago  ingress-nginx  4.11.0
kubernetes  default     1            2023-05-12 12:23:12 1d ago  kubernetes    23.0.0
grafana    default     1            2023-05-12 12:23:12 1d ago  grafana        8.3.0
prometheus-community default     1            2023-05-12 12:23:12 1d ago  prometheus-community 2.57.0
● josip@G14:~$ 
○ △ 0  k3d-k3s-default  default
```

this should be your output.

5. helm install prometheus prometheus-community/kube-prometheus-stack

```
● josip@G14:~$ helm install prometheus prometheus-community/kube-prometheus-stack
NAME: prometheus
LAST DEPLOYED: Fri May 12 12:23:12 2023
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace default get pods -l "release=prometheus"

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Prometheus instances using the operator.
● josip@G14:~$ 
```

6. helm install loki grafana/loki-stack

7. helm list

```
● josip@G14:~$ helm install loki grafana/loki-stack
NAME: loki
LAST DEPLOYED: Fri May 12 12:25:25 2023
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
The Loki stack has been deployed to your cluster. Loki can now be added as a datasource in Grafana.

See http://docs.grafana.org/features/datasources/loki/ for more detail.
● josip@G14:~$ helm list
NAME      NAMESPACE   REVISION      UPDATED                         STATUS    CHART          APP VERS
loki      default     1            2023-05-12 12:25:25.754099674 +0200 CEST  deployed  loki-stack-2.9.10  v2.6.1
grafana   default     1            2023-05-12 12:23:12.866842416 +0200 CEST  deployed  kube-prometheus-stack-45.27.2  v0.65.1
● josip@G14:~$ 
```

You should have a cluster like this in the default namespace

```
● josip@G14:~$ kubectl get pods
NAME                                         READY   STATUS    RESTARTS   AGE
nginx-76d6c9b8c-5bnzj                        1/1    Running   5 (3h25m ago)  27d
tektondemo-may-12-hello-pod                   0/1    Completed  0          111m
tektondemo-may-12-goodbye-pod                 0/1    Completed  0          111m
prometheus-prometheus-node-exporter-hhtvg     1/1    Running   0          3m42s
prometheus-prometheus-node-exporter-b78jn      1/1    Running   0          3m42s
prometheus-prometheus-node-exporter-4q2wf      1/1    Running   0          3m42s
prometheus-kube-state-metrics-858496487b-nxvpn  1/1    Running   0          3m42s
prometheus-kube-prometheus-operator-5cbd7f4cc-ghn52 1/1    Running   0          3m42s
alertmanager-prometheus-kube-prometheus-alertmanager-0 2/2    Running   0          3m32s
grafana-5f5fbddd9-bfnnc                      3/3    Running   0          3m42s
prometheus-prometheus-kube-prometheus-prometheus-0 2/2    Running   0          3m32s
loki-promtail-665vv                           1/1    Running   0          102s
loki-promtail-g6w89                           1/1    Running   0          102s
loki-promtail-8vdks                          1/1    Running   0          102s
loki-0                                         1/1    Running   0          102s
● josip@G14:~$ 
○ △ 0  k3d-k3s-default  default
```

with Prometheus, grafana and loki installed.

8. Next part is to connect to Grafana and insert the databases, we start with the command: `kubectl get svc -n default`

```
● josip@G14:~$ kubectl get svc -n default
NAME           CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes     ClusterIP    <none>        443/TCP         27d
nginx          ClusterIP    10.43.62.169  <none>        80/TCP          27d
prometheus-kube-prometheus-operator  ClusterIP    10.43.165.67  <none>        443/TCP         31m
prometheus-kube-state-metrics       ClusterIP    10.43.68.63   <none>        8080/TCP        31m
prometheus-grafana                 ClusterIP    10.43.22.216  <none>        80/TCP          31m
prometheus-prometheus-node-exporter ClusterIP    10.43.108.20  <none>        9100/TCP        31m
prometheus-kube-prometheus-prometheus ClusterIP  10.43.130.169 <none>        9090/TCP        31m
prometheus-kube-prometheus-alertmanager ClusterIP  10.43.131.38  <none>        9093/TCP        31m
alertmanager-operated               ClusterIP    None          <none>        9093/TCP,9094/TCP,9094/UDP 31m
prometheus-operated                ClusterIP    None          <none>        9090/TCP        31m
loki-memberlist                   ClusterIP    None          <none>        7946/TCP        29m
loki-headless                     ClusterIP    None          <none>        3100/TCP        29m
loki                           ClusterIP  10.43.62.169  <none>        3100/TCP        29m
● josip@G14:~$
```

We need to connect to the svc of grafana which is called prometheus-grafana.

9. To access the grafana GUI we need the secret password from `kubectl get secret`, run this command: `kubectl get secret prometheus-grafana -o jsonpath='{.data.admin-password}' | base64 --decode`

```
● josip@G14:~$ kubectl get secret prometheus-grafana -o jsonpath='{.data.admin-password}' | base64 --decode
● prom-operatorjosip@G14:~$
```

in our case its prom-operator. And the username is of course admin.

10. Now we port-forward into the svc locally with: `kubectl port-forward svc/prometheus-grafana 8082:80`. We selected 8082 here because it's an available port, and type in `localhost:8082` in our browser.
11. Now we can add prometheus and loki as data sources.

Welcome to Grafana

Need help? [Documentation](#) [Tutorials](#) [Community](#) [Public Slack](#)

**Basic**

The steps below will guide you to quickly finish setting up your Grafana installation.

**TUTORIAL DATA SOURCE AND DASHBOARDS**

**Grafana fundamentals**

Set up and understand Grafana if you have no prior experience. This tutorial guides you through the entire process and covers the "Data source" and "Dashboards" steps to the right.

**COMPLETE**

Add your first data source

**COMPLETE**

Create your first dashboard

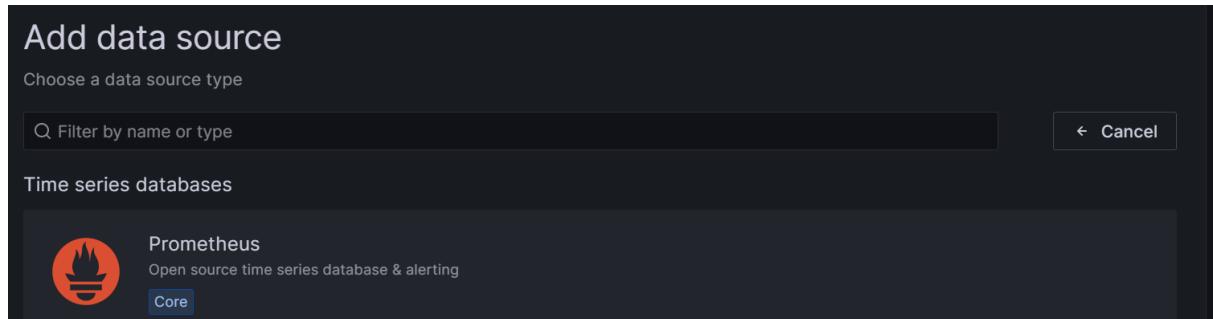
Remove this panel

Dashboards

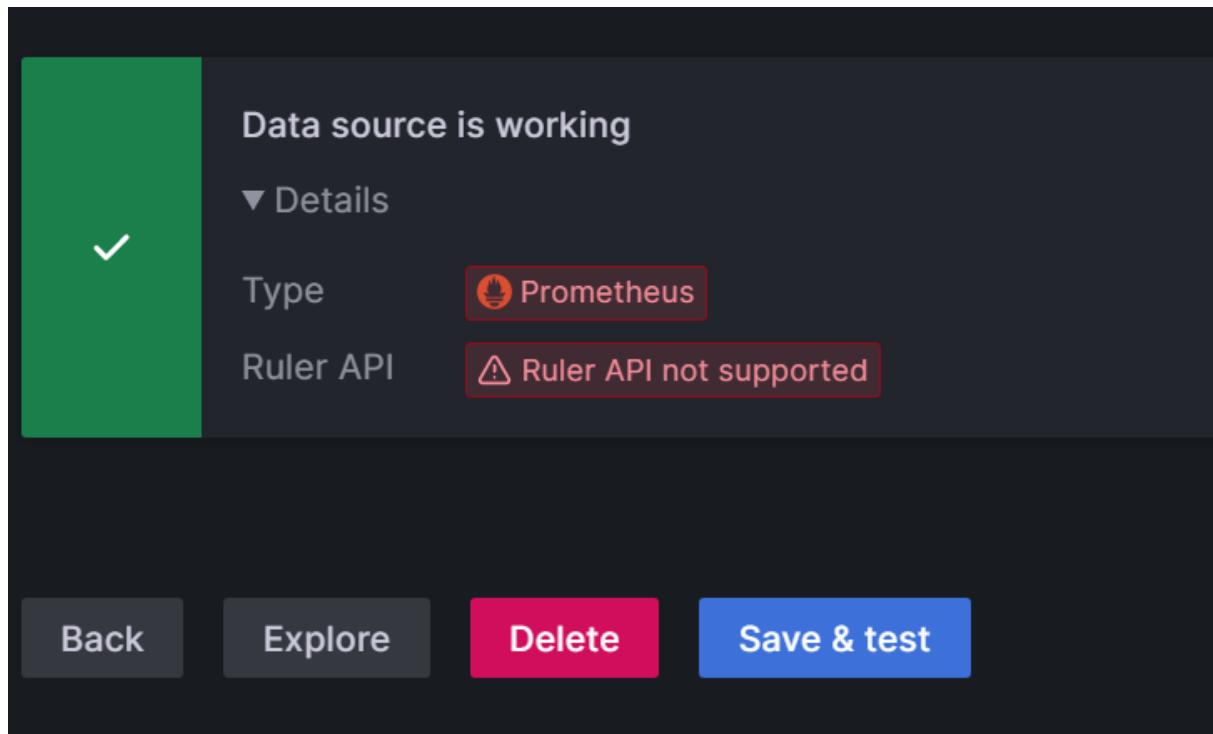
Latest from the blog

12. click on Add your first data source.

13. We select prometheus to begin with.



14. Here we select the url for prometheus. in this case we need the service name of prometheus,namespace and port:  
`http://prometheus-kube-prometheus-prometheus.default.svc.cluster.local:9090`  
this is the url in our case. If it worked you should get a data source with a working message.



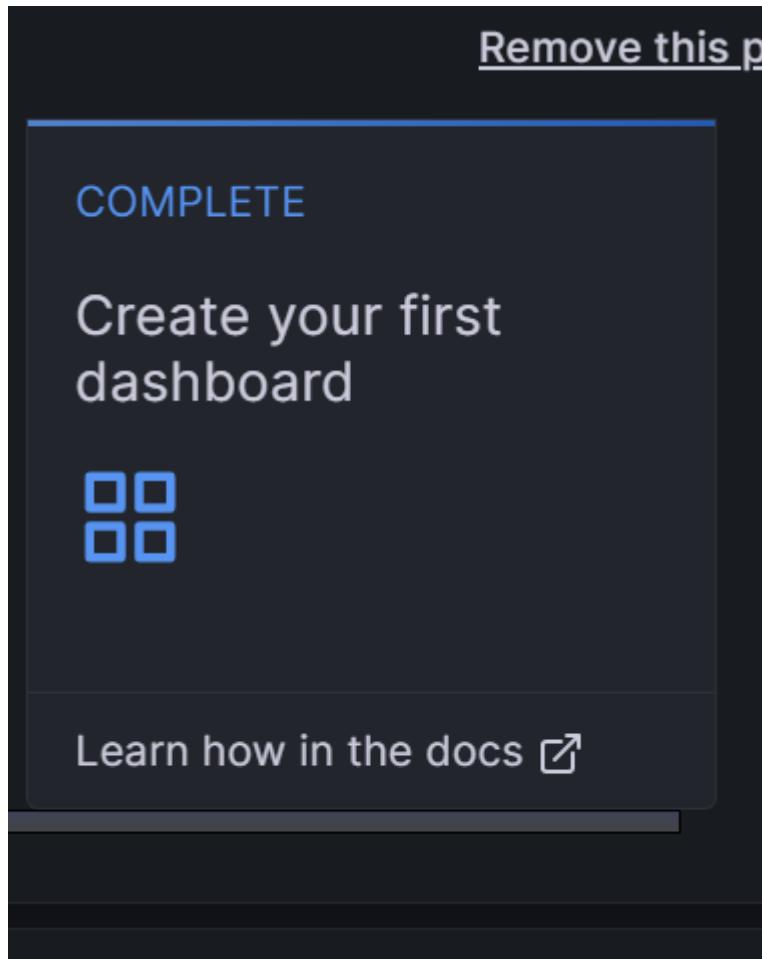
15. For the Loki datasource we almost do the same commands.

The screenshot shows the 'Settings' tab for a Loki data source in Grafana. The configuration section is titled 'Configure your Loki data source below'. It includes an information icon and a note about skipping effort to get Loki as a fully-managed service. Below this, there's a green button labeled '(?) Alerting supported'. The data source is named 'Loki' and is set as the 'Default'. Under the 'HTTP' section, the URL is set to 'http://loki.default.svc.cluster.local:3100'. There are fields for 'Allowed cookies' and 'Timeout'. A large 'Add' button is visible at the bottom left of the configuration area.

16. http://loki.default.svc.cluster.local:3100

The screenshot shows a success message from Grafana: 'Data source connected and labels found.' This message is displayed in a green box with a checkmark icon. Above this message, there is a note about derived fields and a large 'Add' button.

17. To create our first dashboard and to find some info about the cluster. Go to create your first dashboard



from here add visualization.

A screenshot of the Grafana dashboard editor. The main area shows a single panel titled 'Panel Title' with the message 'No data'. Below this, there are tabs for 'Query' (selected), 'Transform', and 'Alert'. The 'Data source' dropdown is set to 'Prometheus'. On the right side, there is a sidebar with various configuration options: 'Search options' (set to 'All'), 'Panel options' (with 'Title' set to 'Panel Title' and 'Description' empty), 'Transparent background' (disabled), 'Panel links', 'Repeat options', and 'Tooltip' (with 'Tooltip mode' empty). At the bottom of the sidebar, there are buttons for 'Run queries', 'Builder', and 'Code'.

add Prometheus as a data source go to the metrics section and here you can enter in a specific query about your cluster.

Metric      Label filters

Select metric    Select label    =    Select value    X    +

+ Operations

or use the predefined ones.

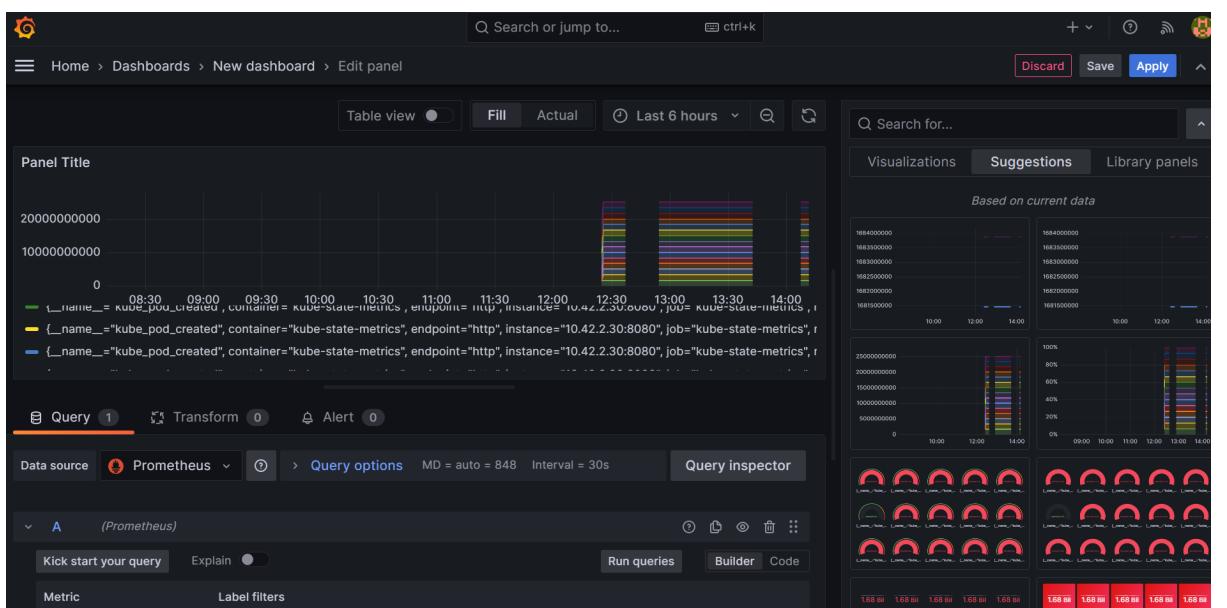
Metric      Label filters

kube\_pod\_created    namespace    =    default    X    +

+ Operations

`kube_pod_created{namespace="default"}`

> Options   Legend: Auto   Format: Time series   Step: auto   Type: Range   Exemplars: false



select a dashboard and save and apply it as easy as that.

## Installing MongoDB and MongoDB Express

For installation follow the steps:

1. git clone <https://github.com/JOJO840/MongoDB-and-Express.git>, it should be looking like this now.

```
josip@G14:~/examensProject/MongoDB-and-Express$ ls
README.md  ingress.yaml  k8s-commands.md  mongo-configmap.yaml  mongo-express.yaml  mongo-secret.yaml  mongo.yaml
josip@G14:~/examensProject/MongoDB-and-Express$
```

2. You need to apply these yaml files in order from the k8s-commands.md file that you have git cloned from the github repo

```
⎈ kubectl apply commands in order

kubectl apply -f mongo-secret.yaml
kubectl apply -f mongo.yaml
kubectl apply -f mongo-configmap.yaml
kubectl apply -f mongo-express.yaml
```

3. Your cluster should be looking like this now with the mongodb and express svcs up and running.

```
josip@G14:~/examensProject/MongoDB-and-Express$ kubectl get all | grep mongo
pod/mongodb-deployment-844789cd64-8qh6w           1/1   Running   0          4m30s
pod/mongo-express-5bf4b56f47-njhsx                1/1   Running   0          2m54s
service/mongodb-service                           ClusterIP   10.43.40.255  <none>           27017/TCP
                                         4m30s
service/mongo-express-service                     LoadBalancer 10.43.97.227    172.20.0.3,172.20.0.4,172.20.0.5  8081:30000/TCP
                                         2m54s
deployment.apps/mongodb-deployment               1/1     1        1          4m30s
deployment.apps/mongo-express                   1/1     1        1          2m54s
replicaset.apps/mongodb-deployment-844789cd64    1         1        1          4m30s
replicaset.apps/mongo-express-5bf4b56f47         1         1        1          2m54s
```

4. To connect to mongodb and express we port-forward into the svc of express, kubectl port-forward svc/mongo-express-service 8082:8081. In this case we connect to port 8082 as 8081 is taken by another app. Then type in localhost:8082 in your browser.

Databases

	Database Name	+ Create Database
	admin	Del
	config	Del
	local	Del

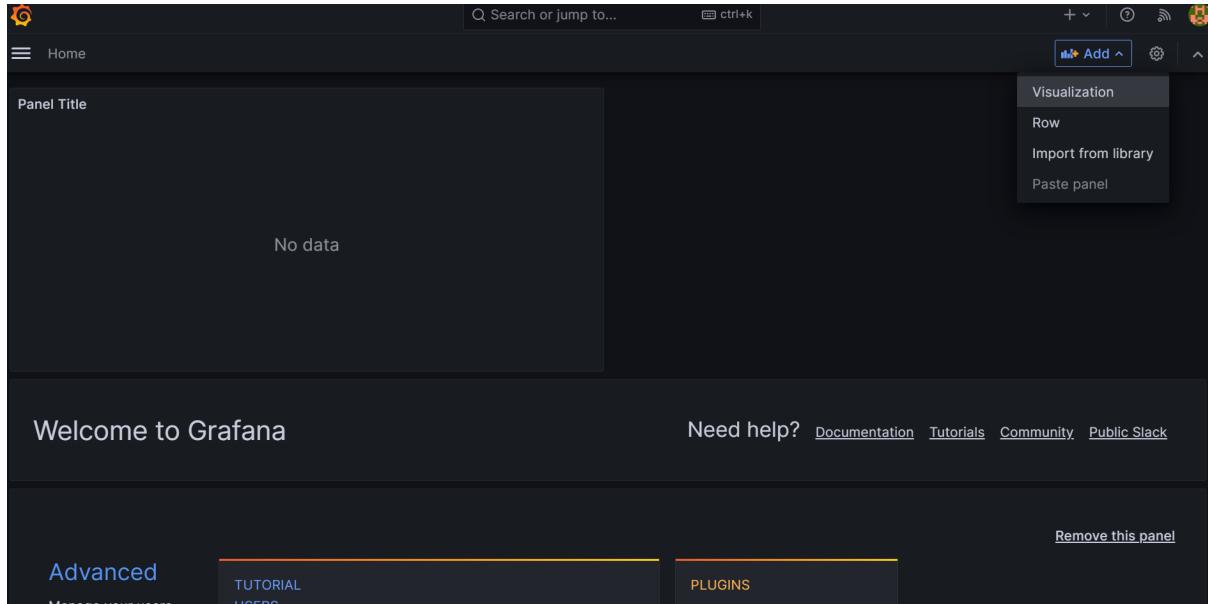
Server Status

Turn on admin in config.js to view server stats!

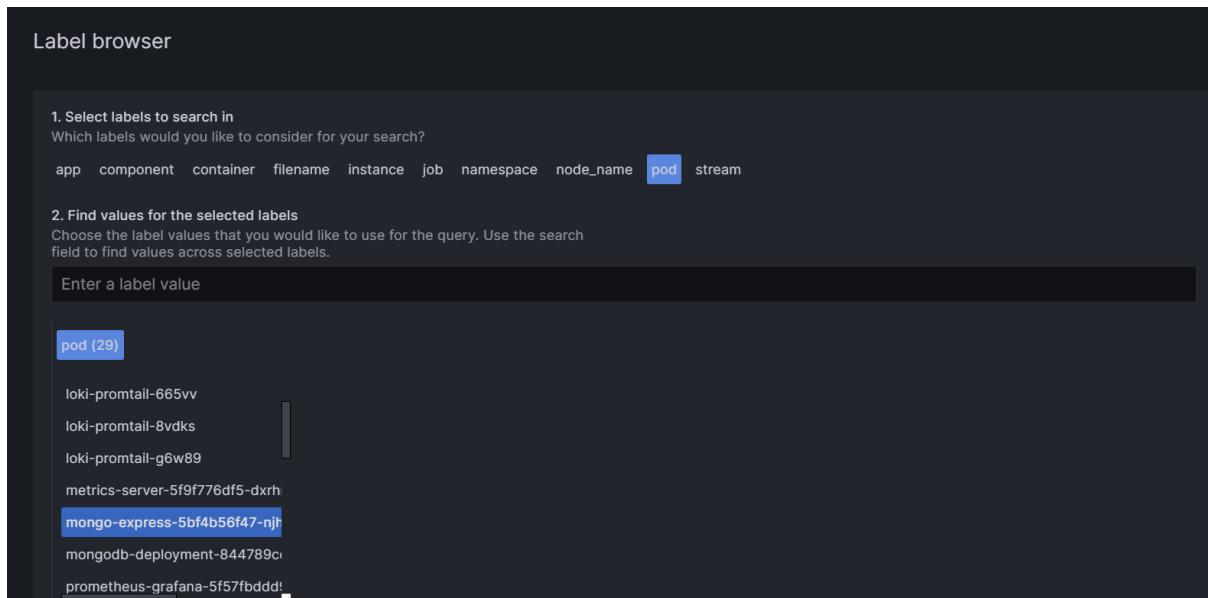
now you're inside of mongodb and express used as a GUI.

## Show logs from Mongodb to Grafana

To show the logs from mongodb in a grafana dashboard we are going to use Loki queries. We start with making a dashboard and choose loki as a datasource. The first step is to go to the home page and press the little add + button, choose visualization.



Here we can choose loki as a datasource and choose the label browser and choose pod and enter in the mongo-express pod from your cluster.



Next we are going into the mongodb-express service with port-forwarding, we are just creating a database here so we can show some logs in grafana.

Databases

	View	Database Name	Project	+ Create Database
	View	admin		Del
	View	config		Del
	View	local		Del

## Server Status

Turn on admin in config.js to view server stats!

Enter in a name for our database and press the Create Database button.

Databases

	View	Database Name	Project	+ Create Database
	View	admin		Del
	View	config		Del
	View	local		Del

## Server Status

Turn on admin in config.js to view server stats!

Now we can see the logs for the database created in grafana. Then click on Save Dashboard. You will have a dashboard in the home section.

Panel Title

```

[+] GET /public/css/bootstrap-theme.min.css 304 1.029 ms - -
[+] GET /public/css/style.css 304 2.509 ms - -
[>] GET /public/css/bootstrap.min.css 304 0.428 ms - -
[>] GET /db/Project/delete_me 200 46.793 ms - 16672

```

Fields

app	mongo-express
container	mongo-express
filename	/var/log/pods/default_mongo-express-5bf4b56f47-njhsx_6877ce93-355e-437c-

Logs

Search options

Panel options

Logs

Time

Unique labels

Common labels

Wrap lines

Prettify JSON

Enable log details

Deduplication

None Exact Numbers Signal

Query 1 Transform 0

Data source http://loki.default:3100 Query options MD = auto = 850 Interval = 30s Query inspector

A ((http://loki.default:3100)) Explain query Run queries Builder Code

Label browser Label filters namespace = default pod = mongo-express-5bf4b56f47-njhsx\_6877ce93-355e-437c-

The screenshot shows a Grafana dashboard interface. At the top, there's a header with a search bar, a 'ctrl+K' hotkey indicator, and various navigation icons. Below the header is a breadcrumb navigation path: Home > Dashboards > Home. To the right of the path are 'Add' and 'Edit' buttons.

The main content area contains a panel titled "Panel Title". Inside the panel, there is a terminal-like window displaying a list of log entries. The log entries are as follows:

```
|> GET /public/img/gears.gif 304 0.715 ms - -
|> GET /public/img/mongo-express-logo.png 304 0.734 ms - -
|> GET /public/collection-fe2ea99c2dc014cd478..min.js 304 0.845 ms - -
|> GET /public/css/theme/rubyblue.css 304 1.619 ms - -
|> GET /public/codemirror-cd38c6e59d64ef16b149..min.js 304 1.502 ms - -
|> GET /public/vendor-d1b820f8a9cf3d5a8c6a.min.js 304 1.418 ms - -
|> GET /public/css/codemirror.css 304 1.399 ms - -
|> GET /public/css/bootstrap-theme.min.css 304 1.353 ms - -
|> GET /public/css/style.css 304 1.398 ms - -
|> GET /public/css/bootstrap.min.css 304 0.339 ms - -
|> GET /db/Project/delete_me 200 33.417 ms - 19124
|> POST /db/Project/delete_me 302 11.964 ms - 86
```

Below the terminal window, the text "Welcome to Grafana" is displayed. To the right of this text is a "Need help?" section containing links to Documentation, Tutorials, Community, and Public Slides. At the bottom left of the panel is a "Advanced" button. On the far right of the dashboard, there is a "Remove this panel" link.

## **6. Discussion:**

The objective of this study was to utilize a bundle of programs from the K8s ecosystem and document the installation process for each component. The tutorial presented in the "Results" section successfully achieved this objective. The clear and functional installation instructions provided by the vendors, along with the readily available official documentation, facilitated straightforward procedures that can be easily replicated. This stable workflow can be credited to the strong community of developers that contribute to the K8s ecosystem regularly.

Additionally, a second tutorial was created to provide an overview of K8s architecture and its components. This tutorial is included in the "Attachment" section. Given the abundance of online resources available about K8s components, we deemed it valuable to include an overview to enhance the proficiency of future trainees.

Regarding the future trends within the industry, Tekton is an example of a tool that is likely to be replaced by other options addressing CI/CD pipelines. Emerging trends are enabling pipelines that can be easily relocated to other platforms. An example of this new approach is GitHub Actions, which, although not currently fully transferable between platforms, has gained popularity as a pipeline solution. When it comes to monitoring capabilities, Grafana remains the preferred choice, offering open-source, cost-effective, and user-friendly features. In terms of cluster setup, K3D stands out as one of the fastest options, offering multi-node capability unlike Minikube, which is limited to a single node

## 7. Conclusion

In conclusion, this study aimed to explore the setup and maintenance of lightweight Kubernetes (K8s) clusters with monitoring and CI/CD pipeline capabilities. The investigation focused on a specific K8s bundle comprising Tekton, Grafana, K3D, and MongoDB, which were selected based on their lightweight nature, ease of use, and compatibility with K8s.

Through interviews with Robin Morero, our LIA Supervisor at Fabled, valuable insights were gained regarding the software selection criteria and the rationale behind choosing these specific tools. The interviews revealed that Fabled prioritizes ease of maintenance, seamless integration, and the ability to upgrade and replace components. The lightweight nature of the chosen software was seen as advantageous, ensuring efficient resource utilization and cost-effectiveness.

The tutorial provided in this study offers step-by-step instructions for installing the K8s bundle and serves as a valuable resource for beginners seeking to leverage K8s effectively. It is important to note that the installation of prerequisites such as WSL 2, Docker Desktop, and access to a Linux CLI is essential but beyond the immediate scope of this research.

Furthermore, the study acknowledges certain limitations, including the level of coverage of K8s architecture and the focus on providing a fundamental understanding rather than going into specific details. The primary goal of this research is to equip individuals with the necessary knowledge and skills to utilize K8s effectively, particularly for lightweight and cost-effective cluster management.

The insights gained from the interviews highlight the importance of managing complexity within K8s clusters. While one approach involves setting up a separate platform team and utilizing vendors for comprehensive management, the preferred approach at Fabled is to minimize complexity by keeping the cluster lightweight and using open source components. This approach reduces maintenance work and provides developers with a seamless experience, aligning with Fabled's emphasis on cost-effectiveness and ease of use.

In summary, this study has provided insights into the setup and maintenance of lightweight K8s clusters with monitoring and CI/CD pipeline capabilities. By exploring the specific K8s bundle, interviews with industry professionals, and the tutorial for installation, this research equips individuals with practical knowledge and resources for harnessing the power of K8s effectively in their respective projects.

**8. Attachments:**

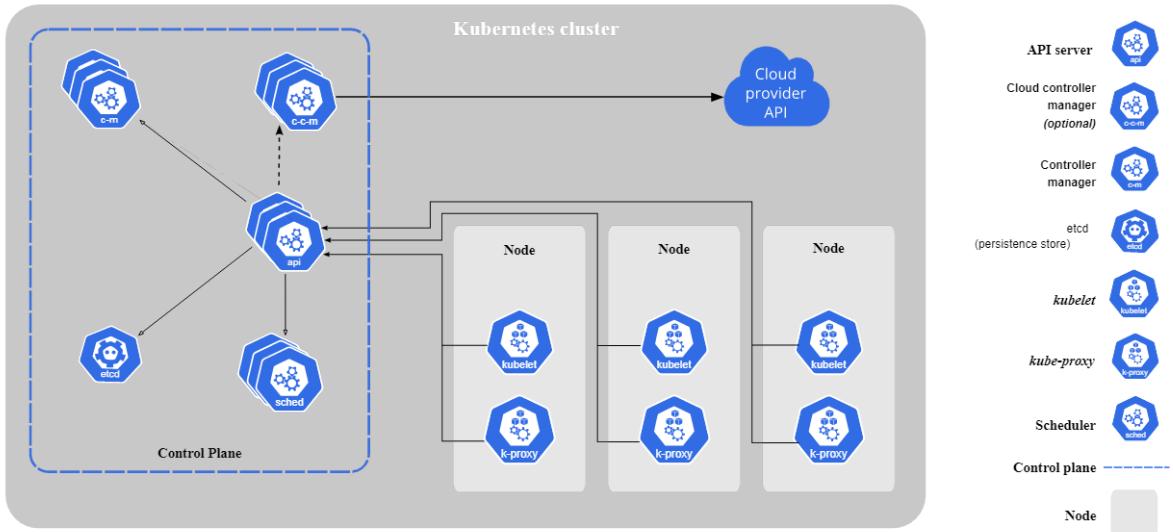
# Kubernetes: An Overview

# k8s Components

## KUBERNETES KOMPONENTER

- ▶ Komponenter
  - ▶ Pod
  - ▶ Service
  - ▶ Ingress
  - ▶ Volumes
  - ▶ Secrets
  - ▶ ConfigMap
  - ▶ StatefulSet
  - ▶ Deployment

A picture of an Cluster



When you deploy Kubernetes, you get a cluster.

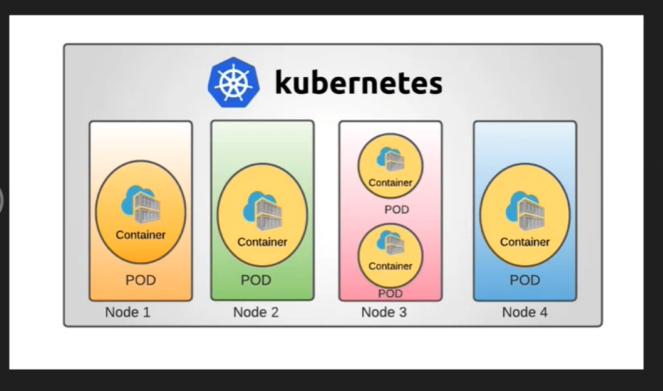
A Kubernetes cluster consists of a set of worker machines, called nodes, that run containerized applications. Every cluster has at least one worker node.

The worker node(s) host the Pods that are the components of the application workload. The control plane manages the worker nodes and the Pods in the cluster. In production environments, the control plane usually runs across multiple computers and a cluster usually runs multiple nodes, providing fault-tolerance and high availability.

### Nodes and Pods:

#### KUBERNETES KOMPONENTER

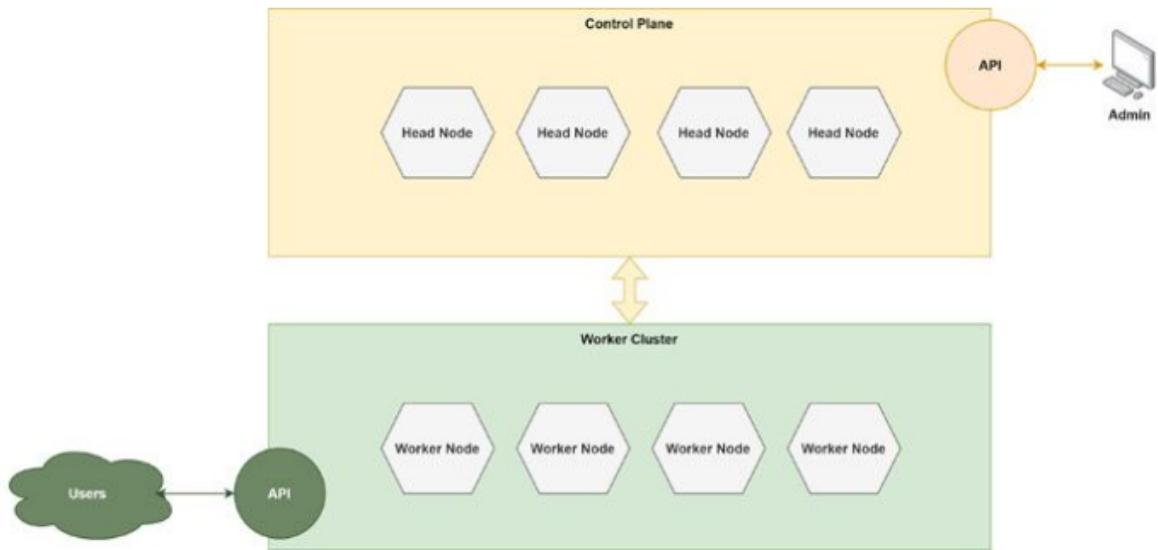
- ▶ Node och Pod
  - ▶ Node (Server eller VM)
  - ▶ Pod (Abstraktion över container)
    - ▶ Vanligtvis 1 container/Pod
    - ▶ Varje Pod => unik IP



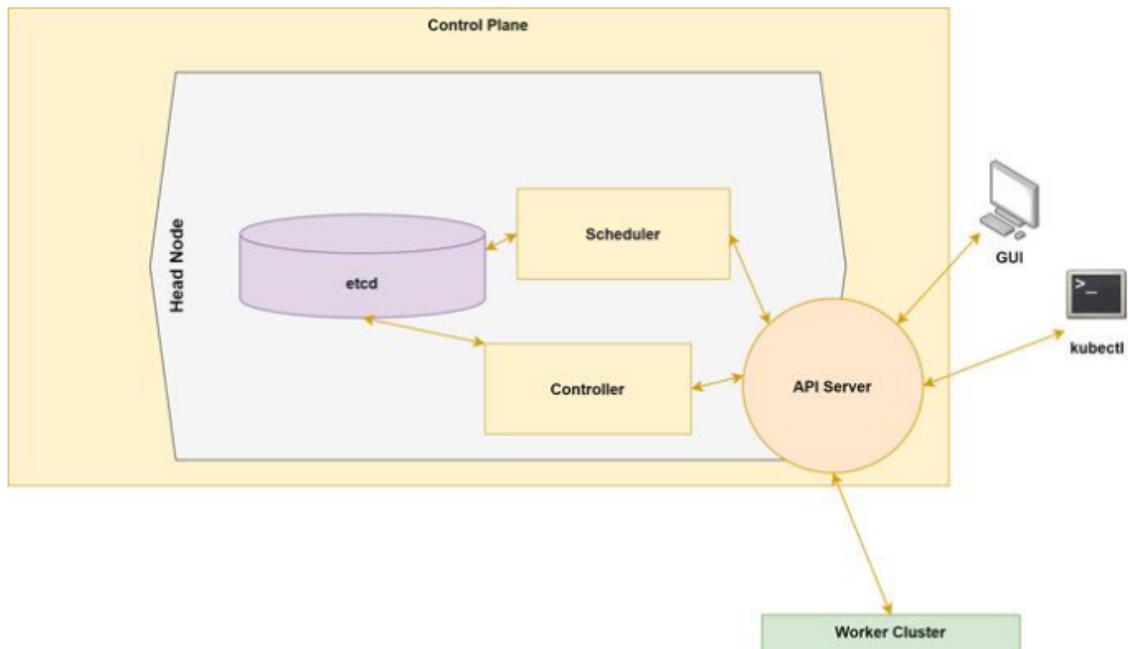
### Master Nodes (head nodes):

A Node is a **VM** or a “**computer**”.

Inside of the worker nodes here is where the containerized apps live, and in the master nodes are the tools to handle that cluster. See the figures.



### A Master Node in detail:



## **etcd (et:si:di)**

Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.

If your Kubernetes cluster uses etcd as its backing store, make sure you have a [back up](#) plan for that data.

You can find in-depth information about etcd in the official [documentation](#).

## **Scheduler**

Control plane component that watches for newly created Pods with no assigned node, and selects a node for them to run on.

Factors taken into account for scheduling decisions include: individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference, and deadlines.

## **Api server**

The core of Kubernetes' control plane is the API server. The API server exposes an HTTP API that lets end users, different parts of your cluster, and external components communicate with one another.

The Kubernetes API lets you query and manipulate the state of API objects in Kubernetes (for example: Pods, Namespaces, ConfigMaps, and Events).

The API server is a component of the Kubernetes control plane that exposes the Kubernetes API. The API server is the front end for the Kubernetes control plane.

Apiserver is designed to scale horizontally—that is, it scales by deploying more instances. You can run several instances of kube-apiserver and balance traffic between those instances.

## **Controller-Manager**

Control plane component that runs controller processes.

The controller is the component that ensures that the cluster remains in the configured state and returns it to that state if it begins to drift. The controller acts as a kind of thermostat that sets a desired state and then strives to maintain it.

Logically, each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process.

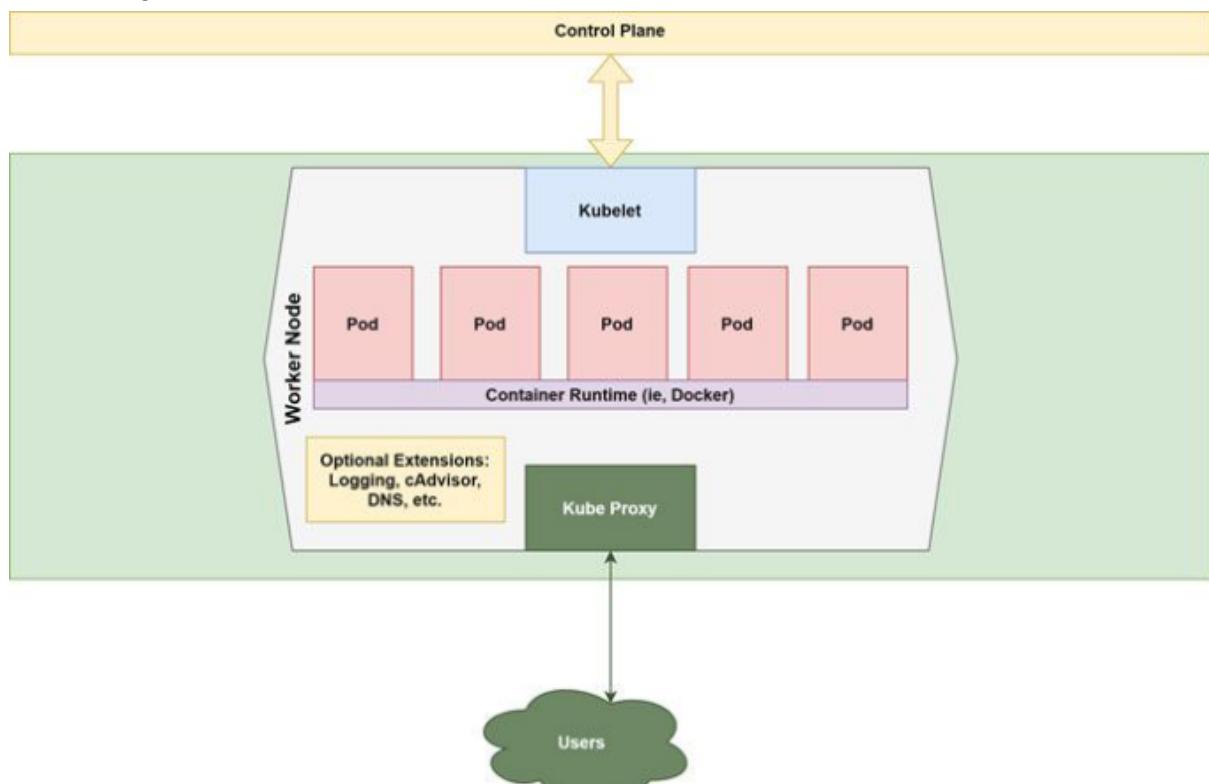
Some types of these controllers are:

- Node controller: Responsible for noticing and responding when nodes go down.
- Job controller: Watches for Job objects that represent one-off tasks, then creates Pods to run those tasks to completion.
- EndpointSlice controller: Populates EndpointSlice objects (to provide a link between Services and Pods).
- ServiceAccount controller: Create default ServiceAccounts for new namespaces.

### Worker Nodes:

The Components inside of a worker node consist of a “Kubelet”, “Kubeproxy”, “Container Runtime” (like Docker for instance) and Pods. Also There are optional Extensions.

See the figure.



## Kubelet

An agent that runs on each node in the cluster. It makes sure that containers are running in a Pod.

The kubelet takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy. The kubelet doesn't manage containers which were not created by Kubernetes.

## Kube-proxy

Kube-proxy is a network proxy that runs on each node in your cluster, implementing part of the Kubernetes Service concept.

[kube-proxy](#) maintains network rules on nodes. These network rules allow network communication to your Pods from network sessions inside or outside of your cluster.

Kube-proxy uses the operating system packet filtering layer if there is one and it's available. Otherwise, kube-proxy forwards the traffic itself.

## Container runtime

The container runtime is the software that is responsible for running containers.

Kubernetes supports container runtimes such as Containerd, CRI-O, and any other implementation of the [Kubernetes CRI \(Container Runtime Interface\)](#).

## Pods

Each pod has a unique IP address, also a pod is an abstraction over a container. Pods are discrete units of work on the node. They are at the replication level. Pods are abstractions that wrap one or more containerized applications.

Pods provide the ability to logically group and isolate containers that run together, while enabling communication between pods on the same machine. The relationship between containers and pods is controlled by Kubernetes deployment descriptors.

Read more on pods on kubernetes doc:

<https://kubernetes.io/docs/concepts/workloads/pods/>

## Deployments and Replica Sets

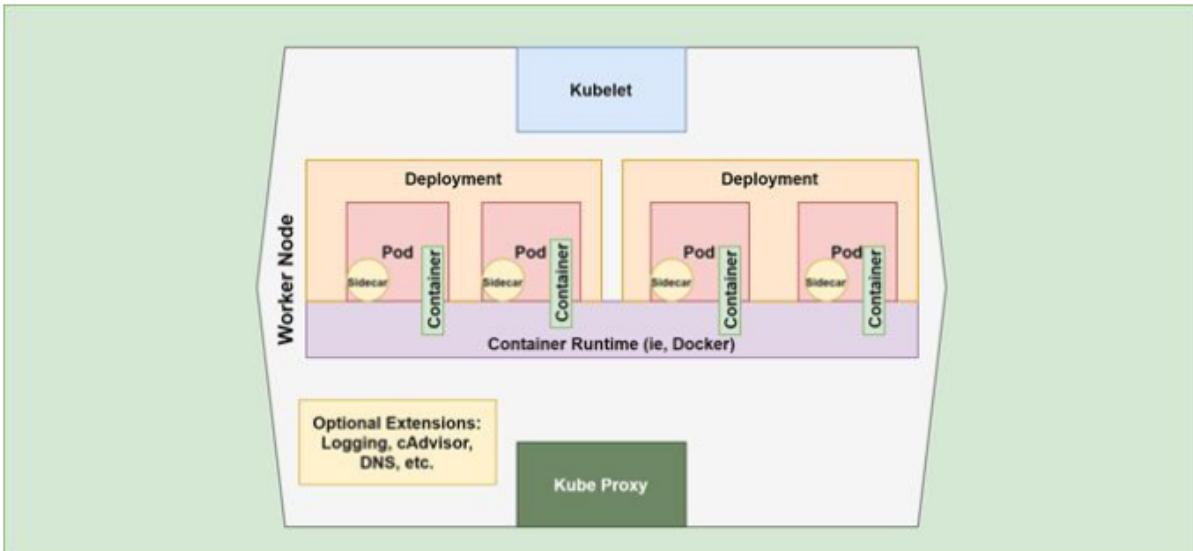
Typically, pods are configured and deployed as parts of ReplicaSets. A ReplicaSet defines the desired execution properties of the pod and makes Kubernetes work to maintain that state.

ReplicaSets are usually defined by a Deployment, which specifies both the parameters of the current ReplicaSet and the strategy to be used (that is, if pods are to be updated or newly created) when managing the cluster.

### Sidecar (Add-ons)

At the pod level, extra functions can be added through add-ons called sidecars. The sidecars handle tasks such as logging at the pod level and collecting statistics.

### A Pod in Detail:



# Service

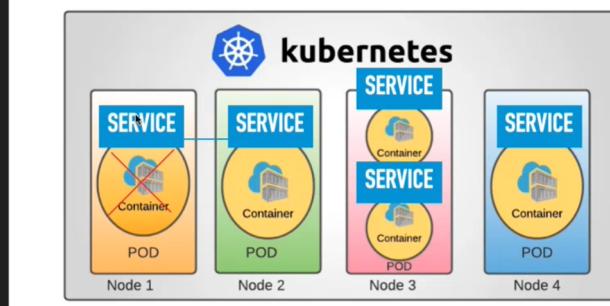
In Kubernetes, a Service is a method for exposing a network application that is running as one or more Pods in your cluster.

A key aim of Services in Kubernetes is that you don't need to modify your existing application to use an unfamiliar service discovery mechanism. You can run code in Pods, whether this is a code designed for a cloud-native world, or an older app you've containerized. You use a Service to make that set of Pods available on the network so that clients can interact with it.

If you use a Deployment to run your app, that Deployment can create and destroy Pods dynamically. From one moment to the next, you don't know how many of those Pods are working and healthy; you might not even know what those healthy Pods are named. Kubernetes Pods are created and destroyed to match the desired state of your cluster. Pods are ephemeral resources (you should not expect that an individual Pod is reliable and durable).

## KUBERNETES KOMPONENTER

- ▶ Service ger Pod fast IP
  - ▶ Om Pod dör lever Service kvar
  - ▶ Podar kan prata ostört!
- ▶ 2 typer
  - ▶ External service (öppen/frontend)
  - ▶ Internal service (privat/backend)



Each Pod gets its own IP address (Kubernetes expects network plugins to ensure this). For a given Deployment in your cluster, the set of Pods running in one moment in time could be different from the set of Pods running that application a moment later.

This leads to a problem: if some set of Pods (call them "backends") provides functionality to other Pods (call them "frontends") inside your cluster, how do the frontends find out and keep track of which IP address to connect to, so that the frontend can use the backend part of the workload?

<https://kubernetes.io/docs/concepts/services-networking/service/>

## Defining a Service

A Service in Kubernetes is an object (the same way that a Pod or a ConfigMap is an object).

You can create, view or modify Service definitions using the Kubernetes API. Usually you use a tool such as `kubectl` to make those API calls for you.

For example, suppose you have a set of Pods that each listen on TCP port 9376 and are labeled as `app.kubernetes.io/name=MyApp`. You can define a Service to publish that TCP listener:

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Applying this manifest creates a new Service named "my-service", which targets TCP port 9376 on any Pod with the `app.kubernetes.io/name: MyApp` label.

Kubernetes assigns this Service an IP address (the *cluster IP*), that is used by the virtual IP address mechanism. For more details on that mechanism, read [Virtual IPs and Service Proxies](#).

## Port definitions

Port definitions in Pods have names, and you can reference these names in the `targetPort` attribute of a Service. For example, we can bind the `targetPort` of the Service to the Pod port in the following way:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app.kubernetes.io/name: proxy
spec:
  containers:
    - name: nginx
      image: nginx:stable
      ports:
        - containerPort: 80
          name: http-web-svc

---
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app.kubernetes.io/name: proxy
  ports:
    - name: name-of-service-port
      protocol: TCP
      port: 80
      targetPort: http-web-svc
```

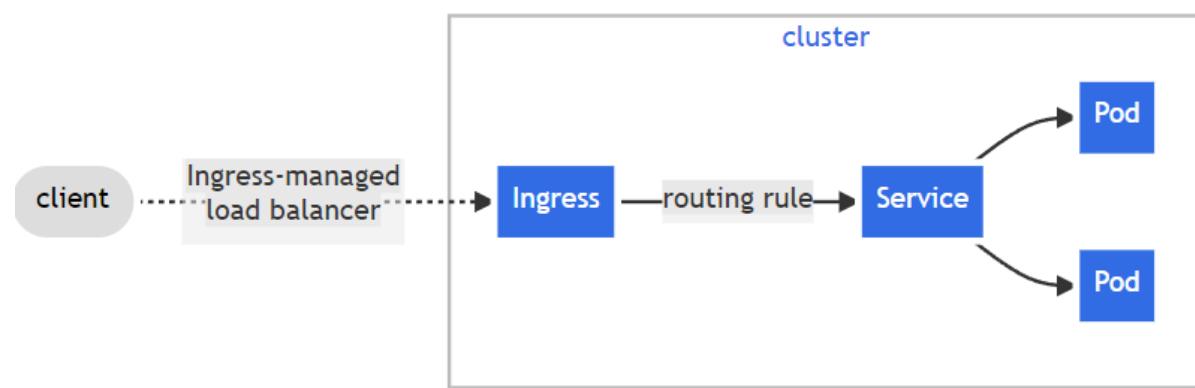
# Ingress

## What is Ingress?

An API object that manages external access to the services in a cluster, typically HTTP.

Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster. Traffic routing is controlled by rules defined on the Ingress resource.

Here is a simple example where an Ingress sends all its traffic to one Service:



An Ingress may be configured to give Services externally-reachable URLs, load balance traffic, terminate SSL / TLS, and offer name-based virtual hosting. An Ingress controller is responsible for fulfilling the Ingress, usually with a load balancer, though it may also configure your edge router or additional frontends to help handle the traffic.

An Ingress does not expose arbitrary ports or protocols. Exposing services other than HTTP and HTTPS to the internet typically uses a service of type [Service.Type=NodePort](#) or [Service.Type=LoadBalancer](#).

Read More on Ingress Here:

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

# Volumes

In Kubernetes, a volume can be thought of as a directory which is accessible to the containers in a pod. We have different types of volumes in Kubernetes and the type defines how the volume is created and its content.

The concept of volume was present with Docker, however the only issue was that the volume was very much limited to a particular pod. As soon as the life of a pod ended, the volume was also lost.

On the other hand, the volumes that are created through Kubernetes is not limited to any container. It supports any or all the containers deployed inside the pod of Kubernetes. A key advantage of Kubernetes volume is, it supports different kind of storage wherein the pod can use multiple of them at the same time.

[https://www.tutorialspoint.com/kubernetes/kubernetes\\_volumes.htm](https://www.tutorialspoint.com/kubernetes/kubernetes_volumes.htm)

On-disk files in a container are ephemeral, which presents some problems for non-trivial applications when running in containers. One problem is the loss of files when a container crashes. The kubelet restarts the container but with a clean state. A second problem occurs when sharing files between containers running together in a [Pod](#). The Kubernetes volume abstraction solves both of these problems. Familiarity with [Pods](#) is suggested.

<https://kubernetes.io/docs/concepts/storage/volumes/>

# Configmap

A ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.

A ConfigMap allows you to decouple environment-specific configuration from your container images, so that your applications are easily portable.

Use a ConfigMap for setting configuration data separately from application code.

For example, imagine that you are developing an application that you can run on your own computer (for development) and in the cloud (to handle real traffic). You write the code to look in an environment variable named `DATABASE_HOST`. Locally, you set that variable to `localhost`. In the cloud, you set it to refer to a Kubernetes Service that exposes the database component to your cluster. This lets you fetch a container image running in the cloud and debug the exact same code locally if needed.

A ConfigMap is not designed to hold large chunks of data. The data stored in a ConfigMap cannot exceed 1 MiB. If you need to store settings that are larger than this limit, you may want to consider mounting a volume or use a separate database or file service.

## ConfigMap object

A ConfigMap is an API [object](#) that lets you store configuration for other objects to use. Unlike most Kubernetes objects that have a `spec`, a ConfigMap has `data` and `binaryData` fields. These fields accept key-value pairs as their values. Both the `data` field and the `binaryData` are optional. The `data` field is designed to contain UTF-8 strings while the `binaryData` field is designed to contain binary data as base64-encoded strings.

The name of a ConfigMap must be a valid [DNS subdomain name](#).

Each key under the `data` or the `binaryData` field must consist of alphanumeric characters, `-`, `_` or `..`. The keys stored in `data` must not overlap with the keys in the `binaryData` field.

Starting from v1.19, you can add an `immutable` field to a ConfigMap definition to create an [immutable ConfigMap](#).

## ConfigMaps and Pods

You can write a Pod `spec` that refers to a ConfigMap and configures the container(s) in that Pod based on the data in the ConfigMap. The Pod and the ConfigMap must be in the same namespace.

**Note:** The `spec` of a static Pod cannot refer to a ConfigMap or any other API objects.

Here's an example ConfigMap that has some keys with single values, and other keys where the value looks like a fragment of a configuration format.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: game-demo
data:
  # property-Like keys; each key maps to a simple value
  player_initial_lives: "3"
  ui_properties_file_name: "user-interface.properties"

  # file-Like keys
  game.properties: |
    enemy.types=aliens,monsters
    player.maximum-lives=5
  user-interface.properties: |
    color.good=purple
    color.bad=yellow
    allow.textmode=true
```

There are four different ways that you can use a ConfigMap to configure a container inside a Pod:

1. Inside a container command and args
2. Environment variables for a container
3. Add a file in read-only volume, for the application to read
4. Write code to run inside the Pod that uses the Kubernetes API to read a ConfigMap

These different methods lend themselves to different ways of modeling the data being consumed. For the first three methods, the kubelet uses the data from the ConfigMap when it launches container(s) for a Pod.

The fourth method means you have to write code to read the ConfigMap and its data. However, because you're using the Kubernetes API directly, your application can subscribe to get updates whenever the ConfigMap changes, and react when that happens. By accessing the Kubernetes API directly, this technique also lets you access a ConfigMap in a different namespace.

Read more on Configmap here:

<https://kubernetes.io/docs/concepts/configuration/configmap/>

## **9. Reference List:**

### **Interview:**

Robin Morero. 2023. Fabled

### **Internet:**

Casey, K. (2022) The Enterprisers Project “Kubernetes by the numbers, in 2022: 11 stats to see”

<https://enterprisersproject.com/article/2022/10/kubernetes-statistics-2022>

(Accessed: 8 May 2023)

CNCF. (2022) “CNCF Annual Survey 2022”.

<https://www.cncf.io/reports/cncf-annual-survey-2022/>

(Accessed: 9 May 2023)

CNCF. (2022) “CNCF Annual Survey 2021”.

<https://www.cncf.io/wp-content/uploads/2022/02/CNCF-Annual-Survey-2021.pdf>

(Accessed: 9 May 2023)

Grafana. 2022. “What is Grafana”.

<https://www.redhat.com/en/topics/data-services/what-is-grafana>

(Accessed: 2 May 2023)

K3D. 2022 “What is k3d”. <https://k3d.io/v5.4.9/>

(Accessed: 1 May 2023)

TechWorld with Nana. 2022, “Microservices explained - the What, Why and How?”.  
[https://www.youtube.com/watch?v=rv4LlmVmVWk&ab\\_channel=TechWorldwithNana](https://www.youtube.com/watch?v=rv4LlmVmVWk&ab_channel=TechWorldwithNana)

(Accessed: 1 May 2023)

TechWorld with Nana 2020 “Complete Application Deployment using Kubernetes Components”.

[https://www.youtube.com/watch?v=EQNO\\_kM96Mo&ab\\_channel=TechWorldwithNana](https://www.youtube.com/watch?v=EQNO_kM96Mo&ab_channel=TechWorldwithNana)

(Accessed: 1 May 2023)

Tekton. 2023. “Overview”.

<https://tekton.dev/docs/concepts/overview/> (Accessed: 1 May 2023)